

CS 344

Operating Systems Lab

Assignment - 2B

Report

Team Members:

1. Braj Nandan Kalya - 190101029
 2. Anurag Singhal - 190123010
 3. Pallav Pandey - 190123044
 4. Alokeveer Mondal - 190123003
-

Task 1

- The scheduler() function in proc.c selects the process for running on the basis of a defined scheduling algorithm.

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;
```

- wakeup1() function fires when a process return from I/O.

```
653 static void
654 wakeup1(void *chan)
655 {
656     struct proc *p;
657
658     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
659         if(p->state == SLEEPING && p->chan == chan) {
660             p->state = RUNNABLE;
661             #ifdef DML
662             p->priority = 3; // relevant for DML - process waited for I/O, and now it's ready to run again
663             #endif
664         }
665 }
666
```

- Allocproc() function in proc.c fires when a new process is created. It allocates resources to the process.

```
//PAGEBREAK: 32
// Look in the process table for an UNUSED proc.
// If found, change state to EMBRYO and initialize
// state required to run in the kernel.
// Otherwise return 0.
static struct proc*
allocproc(void)
{
    struct proc *p;
    char *sp;

    acquire(&ptable.lock);

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
        if(p->state == UNUSED)
            goto found;
```

- Scheduling takes place whenever
 1. In default, time quantum is completed for the current process
 2. In FCFS, the current process is completed.
 3. In SML, either process is completed or a higher priority process is ready.
 4. In DML, either process is completed or a higher priority process is ready or if time quantum is over.

Policy 1: Default Policy

```
451 scheduler(void)
452 {
453     struct proc *p;
454     struct cpu *c = mycpu();
455     c->proc = 0;
456     int index1 = 0;
457     int index2 = 0;
458     int index3 = 0;
459     int x = index1;
460     index1 = x;
461     x = index2;
462     index2 = x;
463     x = index3;
464     index3 = x;
465
466     for(;;){
467         // Enable interrupts on this processor.
468         sti();
469
470         // Loop over process table looking for process to run.
471         acquire(&ptable.lock);
472         #ifdef DEFAULT
473         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
474             if(p->state != RUNNABLE)
475                 continue;
476
477             // Switch to chosen process. It is the process's job
478             // to release ptable.lock and then reacquire it
479             // before jumping back to us.
480             c->proc = p;
481             switchvm(p);
482             p->state = RUNNING;
483
484             switch(&(c->scheduler), p->context);
485             switchkvm();
486
487             // Process is done running for now.
488             // It should have changed its p->state before coming back.
489             c->proc = 0;
490         }
491     }
492     #else
```

Represents the scheduling policy currently implemented at xv6 (with the only difference being the newly defined QUANTA).

Policy 2: First come - First Served

```
493 #ifdef FCFS
494     struct proc *minP = NULL;
495     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
496         if(p->state == RUNNABLE){
497             if (minP!=NULL){
498                 if(p->ctime < minP->ctime)
499                     minP = p;
500             }
501             else
502                 minP = p;
503         }
504     }
505     if (minP!=NULL){
506         p = minP; //the process with the smallest creation time
507         c->proc = p;
508         switchvm(p);
509         p->state = RUNNING;
510         swtch(&c->scheduler, myproc()->context);
511         switchkvm();
512         // Process is done running for now.
513         // It should have changed its p->state before coming back.
514         c->proc = 0;
515     }
516     #else
```

Represents a non-preemptive policy that selects the process with the lowest creation time. The process runs until it no longer needs CPU time (IO / yield / block).

Policy 3: Multi-level queue scheduling

```
361
362 #ifdef SML
363 struct proc* findreadyprocess(int *index1, int *index2, int *index3, uint *priority) {
364     int i;
365     struct proc* proc2;
366     notfound:
367     for (i = 0; i < NPROC; i++) {
368         switch(*priority) {
369             case 1:
370                 proc2 = &ptable.proc[(*index1 + i) % NPROC];
371                 if (proc2->state == RUNNABLE && proc2->priority == *priority) {
372                     *index1 = (*index1 + 1 + i) % NPROC;
373                     return proc2; // found a runnable process with appropriate priority
374                 }
375             case 2:
376                 proc2 = &ptable.proc[(*index2 + i) % NPROC];
377                 if (proc2->state == RUNNABLE && proc2->priority == *priority) {
378                     *index2 = (*index2 + 1 + i) % NPROC;
379                     return proc2; // found a runnable process with appropriate priority
380                 }
381             case 3:
382                 proc2 = &ptable.proc[(*index3 + i) % NPROC];
383                 if (proc2->state == RUNNABLE && proc2->priority == *priority){
384                     *index3 = (*index3 + 1 + i) % NPROC;
385                     return proc2; // found a runnable process with appropriate priority
386                 }
387         }
388     }
389     if (*priority == 1) { //did not find any process on any of the priorities
390         *priority = 3;
391         return 0;
392     }
393     else {
394         *priority -= 1; //will try to find a process at a lower priority
395         goto notfound;
396     }
397     return 0;
398 }
399 #endif
```

Represents a preemptive policy that includes a three priority queues.

This function chooses the process to be run according to their priority.

```
517 #ifdef SML
518 uint priority = 3;
519 p = findreadyprocess(&index1, &index2, &index3, &priority);
520 if (p == 0) {
521     release(&ptable.lock);
522     continue;
523 }
524 c->proc = p;
525 switchvm(p);
526 p->state = RUNNING;
527 swtch(&c->scheduler, myproc()->context);
528 switchkvm();
529 c->proc = 0;
530 #else
```

In this scheduling policy the scheduler will select a process from a lower queue only if no process is ready to run at a higher queue. Moving between priority queues is only available via a system call. Priority 3 is the highest priority.

Policy 4: Dynamic Multi-level queue scheduling

```
401 #ifdef DML
402 /*
403  * this method will find the next process to run
404  */
405 struct proc* findreadyprocess(int *index1, int *index2, int *index3, uint *priority) {
406     int i;
407     struct proc* proc2;
408     notfound:
409     for (i = 0; i < NPROC; i++) {
410         switch(*priority) {
411             case 1:
412                 proc2 = &ptable.proc[(*index1 + i) % NPROC];
413                 if (proc2->state == RUNNABLE && proc2->priority == *priority) {
414                     *index1 = (*index1 + 1 + i) % NPROC;
415                     return proc2; // found a runnable process with appropriate priority
416                 }
417             case 2:
418                 proc2 = &ptable.proc[(*index2 + i) % NPROC];
419                 if (proc2->state == RUNNABLE && proc2->priority == *priority) {
420                     *index2 = (*index2 + 1 + i) % NPROC;
421                     return proc2; // found a runnable process with appropriate priority
422                 }
423             case 3:
424                 proc2 = &ptable.proc[(*index3 + i) % NPROC];
425                 if (proc2->state == RUNNABLE && proc2->priority == *priority){
426                     *index3 = (*index3 + 1 + i) % NPROC;
427                     return proc2; // found a runnable process with appropriate priority
428                 }
429         }
430     }
431     if (*priority == 1) { // did not find any process on any of the priorities
432         *priority = 3;
433         return 0;
434     }
435     else {
436         *priority -= 1; // will try to find a process at a lower priority
437         goto notfound;
438     }
439     return 0;
440 }
441 #endif
```

Represents a preemptive policy similar to Policy 3. The difference is that the process cannot manually change its priority.

Running the full quanta will result in a decrease of priority by 1.

This function chooses the process to be run according to their priority.

```
531 #ifdef DML
532 uint priority = 3;
533 p = findreadyprocess(&index1, &index2, &index3, &priority);
534 if (p == 0) {
535     release(&ptable.lock);
536     continue;
537 }
538 c->proc = p;
539 switchvm(p);
540 p->state = RUNNING;
541 p->tickcounter = 0;
542 swtch(&c->scheduler, myproc()->context);
543 switchkvm();
544 c->proc = 0;
545 #endif
546 #endif
547 #endif
548 #endif
549 release(&ptable.lock);
550
551 }
552 }
```

```

767 void resettickcycle(int *counter) {
768     acquire(&ptable.lock);
769     *counter = 0;
770     release(&ptable.lock);
771 }
772

```

Function created to reset tick counter for each process. It is called when time quantum of a process is completed.

```

737 int set_prio(int priority) {
738     if (priority < 1 || priority > 3)
739         return 1;
740     acquire(&ptable.lock);
741     myproc()->priority = priority;
742     release(&ptable.lock);
743     return 0;
744 }
745

```

This function sets priority of the processes.

```

773 void decpriority(void) {
774     // acquire(&ptable.lock);
775     myproc()->priority = myproc()->priority == 1 ? 1 : myproc()->priority - 1;
776     // release(&ptable.lock);
777 }
778
779 int inctickcounter() {
780     int res;
781     acquire(&ptable.lock);
782     res = ++myproc()->tickcounter;
783     release(&ptable.lock);
784     return res;
785 }

```

- decpriority() decreases the priority of the process by 1 if not already = 1. It is called when the time quantum of a process is over in DML policy.

Following were the changes in makefile.

```
73
74  ifndef SCHEDFLAG
75  SCHEDFLAG := DEFAULT
76  endif
77
78  CC = $(TOOLPREFIX)gcc
79  AS = $(TOOLPREFIX)gas
80  LD = $(TOOLPREFIX)ld
81  OBJCOPY = $(TOOLPREFIX)objcopy
82  OBJDUMP = $(TOOLPREFIX)objdump
83  CFLAGS = -fno-pic -static -fno-builtins
84  CFLAGS += -D $(SCHEDFLAG)
```

```
187  _wc\
188  _zombie\
189  _sanity\
190  _SMLsanity\
191
```

Task 2

yield() and set_prio() system calls were added.

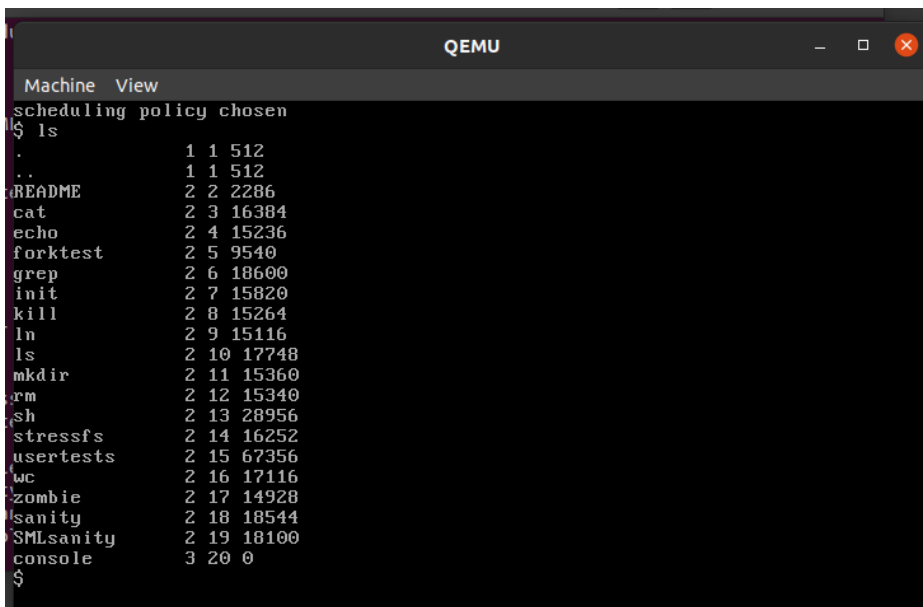
```
23  #define SYS_history 22
24  #define SYS_wait2 23
25  #define SYS_set_prio 24
26  #define SYS_yield 25
```

```
11  #define LOGSIZE (MAXOPB
12  #define NBUF (MAXOPB
13  #define FSSIZE 1000
14  #define QUANTA 5
15
```

```
122  int sys_set_prio(void) {
123      int priority;
124      argint(0, &priority);
125      return set_prio(priority);
126  }
127
128  int sys_yield(void) {
129      yield();
130      return 0;
131  }
```

Task 3

1. Sanity for $n = 3$,
 - sanity for DEFAULT
 - sanity for FCFS
 - sanity for SML
 - sanity for DML
2. SMLsanity with $n = 10$



```
Machine View
scheduling policy chosen
$ ls
.          1 1 512
..         1 1 512
README    2 2 2286
cat        2 3 16384
echo      2 4 15236
forktest  2 5 9540
grep       2 6 18600
init       2 7 15820
kill       2 8 15264
ln         2 9 15116
ls         2 10 17748
mkdir      2 11 15360
rm         2 12 15340
sh         2 13 28956
stressfs   2 14 16252
usertests  2 15 67356
wc         2 16 17116
zombie     2 17 14928
sanity     2 18 18544
SMLsanity  2 19 18100
console    3 20 0
$
```

Sanity and SMLsanity, both are visible in 'ls'.

sanity for DEFAULT with $n = 3$

```
$ sanity 3
S-CPU bound, pid: 5, ready: 2, running: 46, sleeping: 1, turnaround: 49
CPU-bound, pid: 7, ready: 2, running: 49, sleeping: 1, turnaround: 52
S-CPU bound, pid: 8, ready: 2, running: 53, sleeping: 1, turnaround: 56
CPU-bound, pid: 10, ready: 2, running: 57, sleeping: 1, turnaround: 60
S-CPU bound, pid: 11, ready: 3, running: 60, sleeping: 1, turnaround: 64
CPU-bound, pid: 13, ready: 3, running: 66, sleeping: 1, turnaround: 70
I/O bound, pid: 6, ready: 3, running: 67, sleeping: 76, turnaround: 146
I/O bound, pid: 9, ready: 3, running: 70, sleeping: 76, turnaround: 149
I/O bound, pid: 12, ready: 3, running: 75, sleeping: 76, turnaround: 154

CPU bound:
Average ready time: 2
Average running time: 57
Average sleeping time: 1
Average turnaround time: 60

CPU-S bound:
Average ready time: 2
Average running time: 53
Average sleeping time: 1
Average turnaround time: 56

I/O bound:
Average ready time: 3
Average running time: 70
Average sleeping time: 76
Average turnaround time: 149
```

sanity for FCFS with $n = 3$

```
$ sanity 3
CPU-bound, pid: 4, ready: 0, running: 7, sleeping: 0, turnaround: 7
S-CPU bound, pid: 5, ready: 0, running: 11, sleeping: 0, turnaround: 11
CPU-bound, pid: 7, ready: 0, running: 14, sleeping: 0, turnaround: 14
S-CPU bound, pid: 8, ready: 0, running: 16, sleeping: 0, turnaround: 16
CPU-bound, pid: 10, ready: 0, running: 21, sleeping: 0, turnaround: 21
S-CPU bound, pid: 11, ready: 0, running: 23, sleeping: 0, turnaround: 23
I/O bound, pid: 6, ready: 0, running: 26, sleeping: 79, turnaround: 105
I/O bound, pid: 9, ready: 0, running: 28, sleeping: 79, turnaround: 107
I/O bound, pid: 12, ready: 0, running: 33, sleeping: 79, turnaround: 112

CPU bound:
Average ready time: 0
Average running time: 14
Average sleeping time: 0
Average turnaround time: 14

CPU-S bound:
Average ready time: 0
Average running time: 16
Average sleeping time: 0
Average turnaround time: 16

I/O bound:
Average ready time: 0
Average running time: 29
Average sleeping time: 79
Average turnaround time: 108

$
```

sanity for DML with $n = 3$

```
$ sanity 3
CPU-bound, pid: 4, ready: 0, running: 11, sleeping: 1, turnaround: 12
S-CPU bound, pid: 5, ready: 0, running: 16, sleeping: 1, turnaround: 17
CPU-bound, pid: 7, ready: 0, running: 20, sleeping: 1, turnaround: 21
S-CPU bound, pid: 8, ready: 0, running: 27, sleeping: 1, turnaround: 28
CPU-bound, pid: 10, ready: 0, running: 32, sleeping: 1, turnaround: 33
S-CPU bound, pid: 11, ready: 0, running: 36, sleeping: 1, turnaround: 37
I/O bound, pid: 6, ready: 0, running: 38, sleeping: 68, turnaround: 106
I/O bound, pid: 9, ready: 0, running: 40, sleeping: 69, turnaround: 109
I/O bound, pid: 12, ready: 0, running: 43, sleeping: 69, turnaround: 112

CPU bound:
Average ready time: 0
Average running time: 21
Average sleeping time: 1
Average turnaround time: 22

CPU-S bound:
Average ready time: 0
Average running time: 26
Average sleeping time: 1
Average turnaround time: 27

I/O bound:
Average ready time: 0
Average running time: 40
Average sleeping time: 68
Average turnaround time: 108

$
```


sanity for SML with $n = 3$

```
$ sanity 3
S-CPU bound, pid: 5, ready: 2, running: 10, sleeping: 1, turnaround: 13
CPU-bound, pid: 7, ready: 2, running: 13, sleeping: 1, turnaround: 16
S-CPU bound, pid: 8, ready: 2, running: 16, sleeping: 1, turnaround: 19
CPU-bound, pid: 10, ready: 2, running: 17, sleeping: 1, turnaround: 20
S-CPU bound, pid: 11, ready: 2, running: 21, sleeping: 1, turnaround: 24
CPU-bound, pid: 13, ready: 2, running: 24, sleeping: 1, turnaround: 27
I/O bound, pid: 6, ready: 2, running: 29, sleeping: 80, turnaround: 111
I/O bound, pid: 9, ready: 2, running: 34, sleeping: 80, turnaround: 116
I/O bound, pid: 12, ready: 2, running: 38, sleeping: 80, turnaround: 120

CPU bound:
Average ready time: 2
Average running time: 18
Average sleeping time: 1
Average turnaround time: 21

CPU-S bound:
Average ready time: 2
Average running time: 15
Average sleeping time: 1
Average turnaround time: 18

I/O bound:
Average ready time: 2
Average running time: 33
Average sleeping time: 80
Average turnaround time: 115

$
```

SMLsanity with n = 10

```
$ SMLsanity 10
Priority 3, pid: 15, ready: 2, running: 86, sleeping: 82, turnaround: 170
Priority 1, pid: 16, ready: 2, running: 89, sleeping: 82, turnaround: 173
Priority 2, pid: 17, ready: 2, running: 94, sleeping: 82, turnaround: 178
Priority 3, pid: 18, ready: 2, running: 98, sleeping: 82, turnaround: 182
Priority 1, pid: 19, ready: 2, running: 99, sleeping: 82, turnaround: 183
Priority 2, pid: 20, ready: 2, running: 103, sleeping: 82, turnaround: 187
Priority 3, pid: 21, ready: 2, running: 105, sleeping: 82, turnaround: 189
Priority 1, pid: 22, ready: 2, running: 109, sleeping: 82, turnaround: 193
Priority 2, pid: 23, ready: 2, running: 113, sleeping: 82, turnaround: 197
Priority 3, pid: 24, ready: 2, running: 116, sleeping: 82, turnaround: 200
Priority 1, pid: 25, ready: 2, running: 117, sleeping: 82, turnaround: 201
Priority 2, pid: 26, ready: 2, running: 120, sleeping: 82, turnaround: 204
Priority 3, pid: 27, ready: 2, running: 123, sleeping: 82, turnaround: 207
Priority 1, pid: 28, ready: 2, running: 130, sleeping: 82, turnaround: 214
Priority 2, pid: 29, ready: 2, running: 135, sleeping: 82, turnaround: 219
Priority 3, pid: 30, ready: 2, running: 137, sleeping: 82, turnaround: 221
Priority 1, pid: 31, ready: 2, running: 140, sleeping: 82, turnaround: 224
Priority 2, pid: 32, ready: 2, running: 146, sleeping: 82, turnaround: 230
Priority 3, pid: 33, ready: 2, running: 149, sleeping: 82, turnaround: 233
Priority 1, pid: 34, ready: 2, running: 153, sleeping: 82, turnaround: 237
Priority 2, pid: 35, ready: 2, running: 159, sleeping: 82, turnaround: 243
Priority 3, pid: 36, ready: 2, running: 161, sleeping: 82, turnaround: 245
Priority 1, pid: 37, ready: 2, running: 164, sleeping: 82, turnaround: 248
Priority 2, pid: 38, ready: 2, running: 169, sleeping: 82, turnaround: 253
Priority 3, pid: 39, ready: 2, running: 175, sleeping: 82, turnaround: 259
Priority 1, pid: 40, ready: 2, running: 180, sleeping: 82, turnaround: 264
Priority 2, pid: 41, ready: 2, running: 186, sleeping: 82, turnaround: 270
Priority 3, pid: 42, ready: 2, running: 191, sleeping: 82, turnaround: 275
Priority 1, pid: 43, ready: 2, running: 196, sleeping: 82, turnaround: 280
Priority 2, pid: 44, ready: 2, running: 200, sleeping: 82, turnaround: 284

Priority 1:
Average ready time: 2
Average running time: 137
Average sleeping time: 82
Average turnaround time: 221

Priority 2:
Average ready time: 2
Average running time: 142
Average sleeping time: 82
Average turnaround time: 226

Priority 3:
Average ready time: 2
Average running time: 134
Average sleeping time: 82
Average turnaround time: 218

$
```

How to use the patch file:

- Keep xv6 cloned repo and the submitted **patchfile.patch** in the same directory. (Preferably a new folder in desktop to avoid errors)
 - (Link to clone xv6: ***git clone git://github.com/mit-pdos/xv6-public.git***)
- Open this parent directory (that contains patch and xv6) in the terminal and run the following command-
 - ***patch -s -p0 < patchfile.patch***
- The xv6-public folder will now have all the required changes.